

LI ZHENXI

Singapore | zli057@e.ntu.edu.sg

Portfolio: [portfolio-emsghin2iq-as.a.run.app](#) | [linkedin.com/in/zhenxi-li-lzx391](#) | [github.com/zhenxi0901](#)

DevOps & Cloud Engineering Intern | SRE and AI Infrastructure

NTU Computer Science student running production Kubernetes on Google Cloud: migrated a live IoT platform to GKE, own its CI/CD pipeline and cut its cloud bill ~15%. Benchmarks LLM serving and agent workloads on hosted APIs and self-hosted NVIDIA A100 GPUs. ICCV 2025 co-author.

EDUCATION

Nanyang Technological University, Singapore

Jul 2023 - Jun 2027 (expected)

Bachelor of Computer Science | Exchange: Tsinghua University, Luleå University of Technology

SKILLS

Cloud & Infrastructure: Google Cloud Platform (GKE, Cloud Run, Cloud Build, Cloud SQL, IAM) | AWS (EKS, EC2) | Kubernetes | Docker | Terraform | Pulumi | Linux

CI/CD: GitHub Actions | Jenkins | Google Cloud Build | Git | GitHub

Observability & Reliability: Prometheus | Grafana | PromQL | Cloud Monitoring | Alerting | Load testing

AI Infrastructure: LLM serving (SGLang) | LLM benchmarking | A100 GPU | Vertex AI | RAG | MCP | PyTorch

Data & Messaging: Kafka | ZooKeeper | PostgreSQL | Redis | MQTT

Languages: Python | Java | Bash | Go | C/C++

WORK EXPERIENCE

Primustech Pte. Ltd., Singapore | DevOps Engineer Intern

May 2026 - Present

- Migrated a live IoT platform from a single Docker Compose VM to **GKE microservices** (Kafka, Cloud SQL, Redis), provisioned with **Terraform** behind Cloud Armor, TLS and autoscaling.
- Own the **Cloud Build CI/CD pipeline** (63 merged PRs): digest-based promotion, PR and Terraform drift gates, approval-gated releases; moved the repository from Bitbucket to GitHub.
- Cut cloud spend ~15% by tracing a rising monitoring bill to 914 orphaned Kafka topics and shipping an automated cleanup job, plus Cloud SQL and VM right-sizing.
- Built **Grafana dashboards and alerts** for ZooKeeper quorum loss, stalled consumers and crash loops; fixed a ZooKeeper election livelock after node replacement.
- Deployed and operated **AiBE**, the company's agentic AI platform, on **Cloud Run**: multi-agent and RAG services on Vertex AI, with CI/CD, IAM and Secret Manager across dev, staging and production.
- Benchmarked AiBE's LLM stack**: load and latency tests of the agent API on Cloud Run, per-request token cost tracking, and model comparisons on latency and cost.
- Shipped remote HVAC control over MQTT/RPC with rate limits, a kill switch and simulator tests.

ENGIE South East Asia Pte. Ltd., Singapore | Control Systems Intern

May 2025 - Jan 2026

- Configured **BMS-room networks and firewalls** for multi-building projects across sites, linking field DDCs and PLCs back to the central control room.
- Mapped and validated **800+ BACnet/OPC points** from DDCs and PLCs into SCADA across live sites.
- Programmed DDC and Siemens S7-1200 PLC logic and linked site data to a central monitoring platform.

Shanghai Jiao Tong University, China | Research Intern

May 2024 - Aug 2024

- Cut NeRF rendering time **40%** and raised image fidelity 30% through CUDA and PyTorch optimisation.

RESEARCH & PUBLICATIONS

Final Year Project: LLM Agent Workload Benchmarking | NTU HyScale Lab

Aug 2026 - Present

- Contributed a Go benchmark adapter to [ARIES](#), the lab's open-source agent-serving framework, running Toolathlon (109 tasks, 34 MCP servers) in isolated containers.
- Benchmarked 320 LLM agent runs** on a hosted API and self-hosted SGLang: tool calls take 2.7% of agent time, environment setup 23% of wall clock, and 96% of input tokens hit the prefix cache.
- Benchmarked GPU efficiency** of Qwen3-4B on an A100: serving jobs held ~7x the GPU time agents used.
- Load-tested concurrency**: 4 isolated runs per VM gave 3x throughput in 6.4 of 31 GiB memory; root-caused and fixed a Docker race that dropped runs under load.

Publication: [PVChat: Personalized Video Chat with One-Shot Learning](#). ICCV 2025 (co-author).